

READ ME

Q1. (on paper)

Q2. (on paper)

Q3. Three separate sections with the rbf code implemented for different data sets. Each one has a `rbf_reg(xtest,xtrain,ytrain,theta,lamb)` which takes in `x_test`, `x_train`, `y_train`, `lambda`, and `theta` and returns a set of `y_predictions` for the `x_test` set, `rbf_val(xvalid,yvalid,xtrain,ytrain,)` which takes in `x_valid`, `y_valid`, `x_train`, `y_train`, and will run the rbf regression over the given grid of `lambda` and `theta` values then will print the optimal `lambda`, `theta`, and lowest error, and `tst_rmse(y,ytest)` which takes in your `y_predictions` from running `rbf_reg` and will find the error between them and the true `y_test` values. Each of these functions are straightforward to call like any other function.